

CSC411 - Project #4

Yui Chit (Michael) Wong - 999806232

Yijin (Catherine) Wang - 998350476

April 1, 2017

Part 1

Question Explain precisely why the code corresponds to the pseudocode below. Specifically, in your report, explain how all the terms (G_t , π , and the update to θ) are computed, quoting the relevant lines of Python.

REINFORCE, A Monte-Carlo Policy-Gradient Method (episodic)

Input: a differentiable policy parameterization $\pi(a|s, \theta), \forall a \in \mathcal{A}, s \in \mathcal{S}, \theta \in \mathbb{R}^n$

Initialize policy weights θ

Repeat forever:

 Generate an episode $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$, following $\pi(\cdot|\cdot, \theta)$

 For each step of the episode $t = 0, \dots, T - 1$:

$G_t \leftarrow$ return from step t

$\theta \leftarrow \theta + \alpha \gamma^t G_t \nabla_{\theta} \log \pi(A_t|S_t, \theta)$

Pseudocode

Answer As mentioned on the assignment page, the policy function π_{θ} is implemented with a single-hidden-layer of neural network. Since the actions for the bipedal walker is continuous, we have to use a Gaussian distribution on π . Thus we pass the hidden layer into two separately fully connected output, which represents the μ and σ to the normal distribution. The activation function are *tanh* and *softplus* (variation on *ReLU*) respectively. The *sigma* value are also clipped if it is too small or big.

```
1 # 1 layer of hidden unit. Activation is ReLU
2 hidden = fully_connected(
3     inputs=x,
4     num_outputs=hidden_size,
5     activation_fn=tf.nn.relu,
6     weights_initializer=hw_init,
7     weights_regularizer=None,
8     biases_initializer=hb_init,
9     scope='hidden')
10
11 # use last layer of neural network as phi(a, s) (the feature)
12 # mu = phi(s, a)^T dot theta
13 mus = fully_connected(
14     inputs=hidden,
15     num_outputs=output_units,
16     activation_fn=tf.tanh,
17     weights_initializer=mw_init,
18     weights_regularizer=None,
19     biases_initializer=mb_init,
20     scope='mus')
```

```

21
22 # softplus is similar to ReLU. Activation function is  $g(x) = \ln(1+e^x)$ 
23 sigmas = tf.clip_by_value(fully_connected(
24     inputs=hidden,
25     num_outputs=output_units,
26     activation_fn=tf.nn.softplus,
27     weights_initializer=sw_init,
28     weights_regularizer=None,
29     biases_initializer=sb_init,
30     scope='sigmas'),
31     TINY, 5)

```

As for the weight initialization, if there is no weight saved from the previous run, we will initialize the weight θ . There are one w and b for each of the layers (hidden, μ , σ). When initializing the weight to each layer, the program uses *xavierinitialization*, another variation of random weight initialization that keep the scale of the gradients in roughly the same scale.

```

1 # if we have the w's and b's saved, load it. Otherwise initialize it
2 if args.load_model:
3     model = np.load(args.load_model)
4     hw_init = tf.constant_initializer(model['hidden/weights'])
5     hb_init = tf.constant_initializer(model['hidden/biases'])
6     mw_init = tf.constant_initializer(model['mus/weights'])
7     mb_init = tf.constant_initializer(model['mus/biases'])
8     sw_init = tf.constant_initializer(model['sigmas/weights'])
9     sb_init = tf.constant_initializer(model['sigmas/biases'])
10 else:
11     hw_init = weights_init
12     hb_init = relu_init
13     mw_init = weights_init
14     mb_init = relu_init
15     sw_init = weights_init
16     sb_init = relu_init

```

Once everything is initialized, we will start training. For each iteration, we will reset the environment (line 2), then generate the states, actions, and rewards from time 0 to time T . When generating the actions, we will randomly sample from the π normal distribution (line 16). Then based on the *pi.sample()*, we will generate the corresponding action *pi.sample*, and using the action, the new state, reward would be generated. We will keep track of all the states, actions, and rewards in 3 lists (*ep_states*, *ep_actions*, and *ep_rewards*). We will also keep track of the total discounted rewards using the variable G (line 20). Then to obtain G_t , the discounted reward starting from time t , the program calls a culmulation sum function on the *ep_rewards* then subtract it from G (line 30). Thus *returns* would be storing the total discounted rewards for each time from time 0 to $T - 1$.

```

1 for ep in range(16384):
2     obs = env.reset() # reset the environment
3
4     G = 0 # generating all the states and actions and rewards
5     ep_states = []

```

```

6     ep_actions = []
7     ep_rewards = [0]
8     done = False
9     t = 0
10    I = 1
11    while not done:
12        ep_states.append(obs)
13        env.render()
14        # pi.sample() is the list of randomly generated probability
15        # pi_sample becomes the action
16        action = sess.run([pi_sample], feed_dict={x:[obs]})[0][0]
17        ep_actions.append(action)
18        obs, reward, done, info = env.step(action)
19        ep_rewards.append(reward * I)
20        G += reward * I # G is the total discounted reward
21        I *= gamma
22
23        t += 1
24        if t >= MAX_STEPS:
25            break
26    # done generating
27
28    if not args.load_model:
29        # G_t = total - cumulative up to time t.
30        returns = np.array([G - np.cumsum(ep_rewards[:-1])]).T
31        index = ep % MEMORY
32
33        # ep_states contains all the state S_0 to S_{T-1}
34        # ep_actions contains all the actions from A_0 to A_{T-1}
35        # returns (ie reward) contains all the G_t's from t=0 to t=T
36        _ = sess.run([train_op],
37                      feed_dict={x:np.array(ep_states),
38                                   y:np.array(ep_actions),
39                                   Returns:returns })

```

Then we will pass the list of states, actions, and the returns into the training step (line 36-39). Then tensorflow will use the state and the weights to generate a new μ and σ . Then it will compute the log probability of the actions given the generate μ and σ .

```

1 # log probability of y given mu and sigma
2 log_pi = pi.log_prob(y, name='log_pi')

```

The cost function used is $J(\theta) = -\sum [G_t \log \pi(A_t | S_t, \theta)]$. The program uses gradient descent to adjust the θ to minimize the cost function

```

1 # Returns is a 1 x (T-1) array for float (rewards)
2 Returns = tf.placeholder(tf.float32, name='Returns')
3 optimizer = tf.train.GradientDescentOptimizer(alpha)
4 train_op = optimizer.minimize(-1.0 * Returns * log_pi)

```

Part 2

Question Your job is to now write an implementation of REINFORCE that will run for the CartPole-v0 (source code here) environment.

In the Cart Pole task, two actions are possible applying a force pushing left, and applying a force pushing right. Each episode stops when the pole inclides at an angle that larger than a threshold, or when the cart moves out of the frame. The at each time-step before the episode stops, the reward is 1.

The policy function should have two outputs the probability of left and the probability of right. Implement the policy function as a softmax layer (i.e., a linear layer that is then passed through softmax.) Note that a softmax layer is simply a fully-connected layer with a softmax activation.

In your report, detail all the modifications that you had to make to the handout code, one-by-one, and briefly state why and how you made the modifications. Include new code (up to a few lines per modification) that you wrote in your report.

Answer The modifications we made are:

1. We changed the network model.

In part 1, the network has a hidden layer which computes the feature vector of x , and pass the feature vector to one layer that computes μ , and another layer that computes σ , then use the output of μ and σ to generate π and use the π to generate action.

Our network for part 2 is simpler, we don't have a hidden layer. We just have a fully connected layer that takes x as input, and output two output units. Then we use sigmoid as the activation function on the output to make the output in range of $(0, 1)$, and then pass the output from activation function to softmax. The output of softmax is a size 2 vector that each element indicates the probability of an action.

Therefore, we removed the setup of layers we don't need, and built our network as follows. (We also adjusted the value of α and γ .)

```
1 alpha = 1e-6
2 gamma = 0.99
3
4 try:
5     output_units = env.action_space.shape[0]
6 except AttributeError:
7     output_units = env.action_space.n
8
9 input_shape = env.observation_space.shape[0]
10 w = tf.get_variable("w", shape=[input_shape, output_units])
11 b = tf.get_variable("b", shape=[output_units])
12 x = tf.placeholder(tf.float32, shape=(None, input_shape), name='x')
13 y = tf.placeholder(tf.int32, shape=(None, 1), name='y')
14
15 layer1 = tf.sigmoid(tf.matmul(x, w)+b)
16 soft_max = tf.nn.softmax(layer1)
```

2. Since we have bernoulli distribution instead of the normal distribution used in part 1 for π , our π is the output of softmax. Our *pi_sample* gets the action of higher corresponding

probability from softmax. Out \log_pi is the logarithm of softmax. And we use act_pi instead of \log_pi to calculate the cost. We used the line of code which is provided to us to compute act_pi . The code converts y which is a list of actions of shape $n \times 1$ to the one hot encoding matrix which has shape $n \times 2 \times 1$, where n is the number of states. The code also add one more dimension to \log_pi which converts the shape of \log_pi from $n \times 2$ to $n \times 1 \times 2$. The result of multiplication of the two converted matrices is a $n \times 1 \times 1$ matrix which is a list of log probabilities of corresponding actions.

Then we used the act_pi to compute the cost.

```
1 pi_sample = tf.argmax(soft_max , axis=1)
2 log_pi = tf.log(soft_max)
3 act_pi = tf.matmul(tf.expand_dims(log_pi , 1), tf.one_hot(y, 2, axis=1))
4
5 Returns = tf.placeholder(tf.float32 , name='Returns ')
6 optimizer = tf.train.GradientDescentOptimizer(alpha)
7 train_op = optimizer.minimize(-1.0 * Returns * act_pi)
```

The rest of the code are mostly what we got from part 1.


Part 3

Train CartPole using your implementation of REINFORCE. (Use a very small learning rate, and a discount rate $\gamma = 0.99$.) You do not have to train it until it works perfectly (i.e., the episode doesn't stop until time-step 200), but wait at least until the average number of time-steps per episode is 50. (It's better if you wait longer, although you are not required to do that.)

Part 3(a)

Question In your report, include a printout that shows how the weights of the policy function changed and how the average number of time-steps per episode change as you trained.

Answer The change of average number of time-steps per episode over the last 25 episodes are as follows:



Changes of Time-steps per Episode

The change of weights of the policy function are as follows:

Weight at episode 0:

```
[[ 0.55756891 0.44880387 0.45428711 0.01119721]
 [ 0.60239559 -0.53826874 -0.63921934 0.78513044]]
```

Weight at episode 100:

[[0.55751288 0.46140882 0.4542748 -0.00796803]
[0.60192931 -0.55228651 -0.63857907 0.80743229]]
Weight at episode 200:
[[0.55744594 0.47939786 0.45419407 -0.03577379]
[0.60120898 -0.57182312 -0.63765335 0.83893186]]
Weight at episode 300:
[[0.55768639 0.51571321 0.45364195 -0.09334514]
[0.59695876 -0.61137104 -0.63584238 0.90098816]]
Weight at episode 400:
[[0.55669647 0.56195247 0.45266479 -0.17034987]
[0.59164733 -0.66522527 -0.6337384 0.98407555]]
Weight at episode 500:
[[0.55319363 0.61433756 0.45161384 -0.2601867]
[0.58770782 -0.72709823 -0.63237381 1.07760775]]
Weight at episode 600:
[[0.5493471 0.6586284 0.4501721 -0.33945304]
[0.58485538 -0.77985466 -0.63131249 1.1578542]]
Weight at episode 700:
[[0.54611695 0.70909899 0.44816312 -0.4303112]
[0.58224368 -0.83812582 -0.62907767 1.25205159]]
Weight at episode 800:
[[0.54470867 0.74866658 0.44712535 -0.49776825]
[0.58082283 -0.8807537 -0.62635362 1.32611656]]
Weight at episode 900:
[[0.54207498 0.80218232 0.44601309 -0.58896798]
[0.57786518 -0.94014865 -0.62367046 1.42442226]]
Weight at episode 1000:
[[0.53997892 0.85531533 0.44451818 -0.68093479]
[0.57664245 -0.9971112 -0.62065309 1.522493]]
Weight at episode 1100:
[[0.53804773 0.90899348 0.44291493 -0.77372855]
[0.57643706 -1.05087185 -0.61618227 1.62414622]]
Weight at episode 1200:
[[0.53436816 0.98216724 0.44084695 -0.90079987]
[0.57591343 -1.12369418 -0.61135685 1.75792503]]
Weight at episode 1300:
[[0.5304538 1.07761431 0.4385353 -1.06434882]
[0.57597333 -1.21745002 -0.60716885 1.92307317]]
Weight at episode 1400:
[[0.52535737 1.19890761 0.43557847 -1.27195978]
[0.57687956 -1.33788776 -0.6048944 2.12136078]]
Weight at episode 1500:
[[0.51959175 1.30228853 0.4316847 -1.45737672]
[0.57848126 -1.44308734 -0.60325855 2.29336071]]
Weight at episode 1600:
[[0.51401436 1.42097223 0.42739692 -1.66827142]
[0.58055174 -1.56275809 -0.60239846 2.48555636]]
Weight at episode 1700:


```

[[ 0.50750446 1.53061366 0.42203677 -1.87267733]
 [ 0.58143383 -1.6764518 -0.60153365 2.66817904]]
Weight at episode 1800:
[[ 0.500332 1.65737295 0.41609058 -2.10924554]
 [ 0.58377492 -1.80496466 -0.59965515 2.88181019]]
Weight at episode 1900:
[[ 0.49324197 1.77153444 0.41069824 -2.32351661]
 [ 0.58694774 -1.92046285 -0.59812069 3.07449174]]
Weight at episode 2000:
[[ 0.48373222 1.88861907 0.40408787 -2.54988098]
 [ 0.59000146 -2.04181719 -0.59656841 3.27574563]]
Weight at episode 2100:
[[ 0.47449186 2.02448845 0.3969498 -2.81081009]
 [ 0.59316611 -2.1812861 -0.59471154 3.50856233]]
Weight at episode 2200:
[[ 0.46633637 2.13581109 0.3906489 -3.0269537 ]
 [ 0.59572083 -2.29615307 -0.59274048 3.70167685]]
Weight at episode 2300:
[[ 0.45680553 2.26001072 0.38414222 -3.27014399]
 [ 0.59872729 -2.42363763 -0.59078163 3.91966486]]
Weight at episode 2400:
[[ 0.44795224 2.381109 0.37753391 -3.50601459]
 [ 0.60229701 -2.54769492 -0.5885542 4.13105249]]
Weight at episode 2500:
[[ 0.43924668 2.49611187 0.37183547 -3.7278161 ]
 [ 0.60309607 -2.66729498 -0.58754814 4.32853603]]
Weight at episode 2600:
[[ 0.42905468 2.59504628 0.3656635 -3.929003 ]
 [ 0.6061002 -2.77073932 -0.58561844 4.50803375]]
Weight at episode 2700:
[[ 0.41852862 2.70760155 0.35971603 -4.15303564]
 [ 0.60998231 -2.88641882 -0.58367425 4.70958757]]
Weight at episode 2800:
[[ 0.40686017 2.81890082 0.35378745 -4.37579584]
 [ 0.61278868 -3.00272322 -0.58225054 4.90878916]]
Weight at episode 2900:
[[ 0.39934814 2.92078543 0.34896737 -4.57286978]
 [ 0.61418551 -3.10798955 -0.58114028 5.08576536]]
Weight at episode 3000:
[[ 0.39068386 3.01555085 0.34464303 -4.75818443]
 [ 0.61748219 -3.20530534 -0.57989573 5.25346327]]

```

Part 3(b)

Question Explain why the final weights you obtained make sense. This requires understanding the what each input dimension means




Diagram representation of the cartpole system

Answer In order to understand why does the final weight make sense, we need to first understand what each input dimension means. The size of the input dimension is 4, which are [position of cart, velocity of cart, angle of pole, rotation rate (angular velocity) of pole]. In the source code, they are denoted by $[x, x_dot, theta, theta_dot]$ (dot means first derivative). Thus at each iteration when we observe the environment, we will observe the 4 components of the cart. The threshold for x is 2.4 units away from the center, and the threshold for $theta$ is 15 degrees from the vertical.

When looking at the final weights, there are some observations that we made:

- For velocity, angle, and angular velocity, the system tries to balance it using the weights. The weights for those components are approximately the same in magnitude but with different signs.
- For angular velocity, if it is negative (falling to the left), $theta_dot \times w_{theta_dot_L}$ will be positive, while $theta_dot \times w_{theta_dot_R}$ will be negative. Thus it will want the cart to be pushed towards the left side (0) to reduce the angular velocity. On the other hand, if angular velocity is positive (falling to the right), it will want the cart to be pushed towards the right (1) to cancel out the momentum.
- Similarly for velocity, since we also do not want to exceed the 2.4 unit threshold of x , when the position of the cart is negative (to the left of center), $x_dot \times w_{x_dot_L}$ will be negative while $x_dot \times w_{x_dot_R}$ will be positive. Thus based on velocity, the cart will want to be pushed towards the right (1), vice versa.
- Same explanation goes for $theta$. The system would want to use the weight to keep the $theta$ to be approximately 0 (vertically upright).
- The trained system prioritize balancing velocity, angle, and angular velocity over the distance of cart. This can be observed because the system uses the weight to balance out the velocity, angle, and angular velocity (as described in point 1). However, the sign of w_{x_L} and w_{x_R} are the same, and unlike the other components, the magnitude is relatively not close. The potential

reason for this can be because the system is not done training yet. We are only reaching an average step of 50 as supposed to reaching average reward of 195 over 100 consecutive trials.

- One minor observation is that the magnitude of w_{x_dot} and w_{theta_dot} are both bigger than magnitude of w_x and w_{theta} respectively. This make sense because when the system is calculating the new state of the cart given an action, the velocity's are multiplied with $\tau = 0.02$, which is the time interval. Thus even though the magnitude of the weight is big, the resulting component would be smaller.